

Ecommerce

1 Artificial Intelligence Capstone Project on E-Commerce

1.0.1 Project Task: Week 1

1.0.2 Importing libraries and datasets

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

import re
from nltk import word_tokenize
from nltk.tokenize import WordPunctTokenizer
from nltk.stem.porter import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
# import string
import warnings
# ! pip install wordcloud
# from wordcloud import WordCloud

from sklearn.preprocessing import LabelEncoder, LabelBinarizer

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression, RidgeClassifier, \
    SGDClassifier
from sklearn.naive_bayes import MultinomialNB, GaussianNB, BernoulliNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, \
    GradientBoostingClassifier, AdaBoostClassifier, BaggingClassifier
```

```

from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, \
    classification_report, precision_score, recall_score, roc_curve, \
    roc_auc_score, auc

import tensorflow as tf
from tensorflow import keras
from sklearn.utils import class_weight
from sklearn.preprocessing import label_binarize
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D, Dropout, GRU
from keras.models import Sequential
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import RandomizedSearchCV, KFold
from sklearn.preprocessing import MinMaxScaler

```

Using TensorFlow backend.

```

[2]: train = pd.read_csv("train_data.csv")
train.head()

```

```

[2]:

```

	name	brand	\
0	All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...	Amazon	
1	Amazon - Echo Plus w/ Built-In Hub - Silver	Amazon	
2	Amazon Echo Show Alexa-enabled Bluetooth Speak...	Amazon	
3	Fire HD 10 Tablet, 10.1 HD Display, Wi-Fi, 16 ...	Amazon	
4	Brand New Amazon Kindle Fire 16gb 7" Ips Displ...	Amazon	

	categories	\
0	Electronics,iPad & Tablets,All Tablets,Fire Ta...	
1	Amazon Echo,Smart Home,Networking,Home & Tools...	
2	Amazon Echo,Virtual Assistant Speakers,Electro...	
3	eBook Readers,Fire Tablets,Electronics Feature...	
4	Computers/Tablets & Networking,Tablets & eBook...	

	primaryCategories	reviews.date	\
0	Electronics	2016-12-26T00:00:00.000Z	
1	Electronics,Hardware	2018-01-17T00:00:00.000Z	
2	Electronics,Hardware	2017-12-20T00:00:00.000Z	
3	Office Supplies,Electronics	2017-08-04T00:00:00.000Z	
4	Electronics	2017-01-23T00:00:00.000Z	

	reviews.text	\
0	Purchased on Black FridayPros - Great Price (e...	
1	I purchased two Amazon in Echo Plus and two do...	
2	Just an average Alexa option. Does show a few ...	
3	very good product. Exactly what I wanted, and ...	

4 This is the 3rd one I've purchased. I've bough...

```
      reviews.title sentiment
0      Powerful tablet  Positive
1  Amazon Echo Plus AWESOME  Positive
2      Average      Neutral
3      Greattttttt  Positive
4      Very durable!  Positive
```

```
[3]: test_val= pd.read_csv("test_data_hidden.csv")
test_val.head()
```

```
[3]:                                     name  brand \
0  Fire Tablet, 7 Display, Wi-Fi, 16 GB - Include...  Amazon
1  Amazon Echo Show Alexa-enabled Bluetooth Speak...  Amazon
2  All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...  Amazon
3  Brand New Amazon Kindle Fire 16gb 7" Ips Displ...  Amazon
4  Amazon Echo Show Alexa-enabled Bluetooth Speak...  Amazon
```

```
                                     categories  primaryCategories \
0  Fire Tablets,Computers/Tablets & Networking,Ta...  Electronics
1  Computers,Amazon Echo,Virtual Assistant Speake...  Electronics,Hardware
2  Electronics,iPad & Tablets,All Tablets,Fire Ta...  Electronics
3  Computers/Tablets & Networking,Tablets & eBook...  Electronics
4  Computers,Amazon Echo,Virtual Assistant Speake...  Electronics,Hardware
```

```
      reviews.date \
0  2016-05-23T00:00:00.000Z
1  2018-01-02T00:00:00.000Z
2  2017-01-02T00:00:00.000Z
3  2017-03-25T00:00:00.000Z
4  2017-11-15T00:00:00.000Z
```

```
      reviews.text \
0  Amazon kindle fire has a lot of free app and c...
1  The Echo Show is a great addition to the Amazo...
2  Great value from Best Buy. Bought at Christmas...
3  I use mine for email, Facebook ,games and to g...
4  This is a fantastic item & the person I bought...
```

```
      reviews.title sentiment
0      very handy device  Positive
1  Another winner from Amazon  Positive
2  simple to use and reliable so far  Positive
3      Love it!!!  Positive
4      Fantastic!  Positive
```

```
[4]: test= pd.read_csv("test_data.csv")
test.head()
```

```
[4]:
```

	name	brand	\
0	Fire Tablet, 7 Display, Wi-Fi, 16 GB - Include...	Amazon	
1	Amazon Echo Show Alexa-enabled Bluetooth Speak...	Amazon	
2	All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...	Amazon	
3	Brand New Amazon Kindle Fire 16gb 7" Ips Displ...	Amazon	
4	Amazon Echo Show Alexa-enabled Bluetooth Speak...	Amazon	

	categories	primaryCategories	\
0	Fire Tablets,Computers/Tablets & Networking,Ta...	Electronics	
1	Computers,Amazon Echo,Virtual Assistant Speake...	Electronics,Hardware	
2	Electronics,iPad & Tablets,All Tablets,Fire Ta...	Electronics	
3	Computers/Tablets & Networking,Tablets & eBook...	Electronics	
4	Computers,Amazon Echo,Virtual Assistant Speake...	Electronics,Hardware	

	reviews.date	\
0	2016-05-23T00:00:00.000Z	
1	2018-01-02T00:00:00.000Z	
2	2017-01-02T00:00:00.000Z	
3	2017-03-25T00:00:00.000Z	
4	2017-11-15T00:00:00.000Z	

	reviews.text	\
0	Amazon kindle fire has a lot of free app and c...	
1	The Echo Show is a great addition to the Amazo...	
2	Great value from Best Buy. Bought at Christmas...	
3	I use mine for email, Facebook ,games and to g...	
4	This is a fantastic item & the person I bought...	

	reviews.title
0	very handy device
1	Another winner from Amazon
2	simple to use and reliable so far
3	Love it!!!
4	Fantastic!

1.0.3 Exploratory Data Analysis

```
[17]: train.duplicated().sum(), test.duplicated().sum(), test_val.duplicated().sum()
```

```
[17]: (2, 3, 3)
```

Train dataset contains 58 duplicate records and train dataset contains 3 duplicate records.

```
[5]: train = train[train.duplicated()==False]
train.shape
```

```
[5]: (3942, 8)
```

```
[6]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3942 entries, 0 to 3999
Data columns (total 8 columns):
name                3942 non-null object
brand               3942 non-null object
categories           3942 non-null object
primaryCategories    3942 non-null object
reviews.date         3942 non-null object
reviews.text         3942 non-null object
reviews.title        3932 non-null object
sentiment            3942 non-null object
dtypes: object(8)
memory usage: 277.2+ KB
```

```
[7]: test_val.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
name                1000 non-null object
brand               1000 non-null object
categories           1000 non-null object
primaryCategories    1000 non-null object
reviews.date         1000 non-null object
reviews.text         1000 non-null object
reviews.title        997 non-null object
sentiment            1000 non-null object
dtypes: object(8)
memory usage: 62.6+ KB
```

Train dataset contains 10 missing values in 'reviews.title' column and test dataset contains 3 missing values in 'reviews.title' column.

```
[8]: pd.set_option('display.max_colwidth',200)
```

Reviews containing Positive Sentiments

```
[9]: train[train.sentiment=='Positive'][['reviews.text', 'reviews.title']].head(10)
```

```
[9]: reviews.text \
0    Purchased on Black FridayPros - Great Price (even off sale)Very powerful and
```

fast with quad core processors Amazing soundWell builtCons -Amazon ads, Amazon need this to subsidize the tablet and wi...

1 I purchased two Amazon in Echo Plus and two dots plus four fire sticks and the hub Philips hue for lamp for the family at Christmas 2017. I,Ãm so happy with these purchases and learning so much w...

3

very good product. Exactly what I wanted, and a very good price

4 This is the 3rd one I've purchased. I've bought one for all of my nieces. No other case compares to this one. It has held protected the tablet so many times from them dropping it.

5

This is a great product. Light weight. I wish it has wifi to download from online.

7

Purchased this for my son. Has room to upgrade memory to allow more books & games. But the speakers could be better or located in a better position.

8 Bought this for my mom and it was just what she needed and at a great price. Been wanting to get an Ipad for myself, but think this might be a great less expensive option for me as well.

10 I got this tablet to replace my sons old one, I love the adult/child profile and the ability to have the 2 year replacement warranty. The case has also came in handy many times.

11

Great product for the kids gaming apps parental controls to make sure you can monitor kids and prevent unwanted app purchases

12

Love the choice of colors. Have two kindles of my own and purchased this for a gift.

```

                                reviews.title
0                                Powerful tablet
1                                Amazon Echo Plus AWESOME
3                                Greatttttttt
4                                Very durable!
5                                You will love it
7                                Great for kids or smaller needs
8                                Great tablet
10                               Great Tablet
11                               Works great
12 great pad for both children and adults
```

Reviews containing Neutral Sentiments

```
[10]: train[train.sentiment=='Neutral'][['reviews.text','reviews.title']].head(10)
```

```
[10]:                                reviews.text \
2
```

Just an average Alexa option. Does show a few things on screen but still limited.

6 My 7-year old daughter saved up to by this. Her brother bought the 8GB about a year earlier, so new she needed more space. The OS is a bit clunky, and less intuitive then on higher priced tablets,...

17

Not as good as before the old kindle, just seams to work better

59 There is nothing spectacular about this item but also nothing majorly wrong with it. The biggest flaw is that this is geared to kids and there is no way that I have found searching settings or onl...

95

It's unfair for me to rate this product cause I have not even taken it out of the box to set it up.

114

I bought this as s present for my 65 year old grandma. She loves it. Very easy to operate. No issues

146

Bought this tablet for 8 year old. It holding up good & she loves it. She enjoys playing her games & being able to get on the internet.

147 bought a few kindles in the past but this time one of it came defective. the port was bent and it was hard to charge but still possible. comes in 4 different color. was 16gb enough space for kids,...

148

Not a substitute for an iPad, but a really good tablet for reading and minimal internet usage.

187

This device is a good if you are looking for a starter tablet for a young individual.

	reviews.title
2	Average
6	OK For Entry Level Tablet
17	Not as good as before
59	Does what it says, missing one key feature
95	Haven't set it up yet
114	Solid tablet
146	Fire tablet
147	Came defective
148	Good Reader
187	Good for 4 year old

Reviews containing Negative Sentiments

```
[11]: train[train.sentiment=='Negative'][['reviews.text','reviews.title']].head(10)
```

```
[11]: reviews.text \
9
```

was cheap, can not run chrome stuff, returned to store.
 97 Worthless, except as a regular echo and a poor excuse for video chat. I
 love my echo devices, bathroom, pool, kitchen, other places where I may need
 hands free, voice activated music and info. My ...
 104
 Too bad Amazon turned this tablet into a big advertising tool. Many apps dont
 work and the camera is not good.
 121 I bought this Kindle for my 7 year old grand-daughter. I bought a warranty
 for it. I bought it in August, I have already had to replace it. The charger
 connection got loose and was not charging. W...
 150 I am reading positive reviews and wish I could say the same. Best Buy is
 great, so this is not a reflection on them, just our experience with the
 product. We have had this product for just over on...
 151 I have to say it was a little confusing and frustrating when i was not
 getting the verification code from amazon , i waited for 20 minutes then i
 requested another code, nothing... then a nother o...
 249
 It's a good device for children because they don't know any better
 267 the speaker
 voice quality is terrible compare the similar size my logitech UE BOOM.the price
 is too high, even I got on promotion with \$79
 368 Needs to be a stand alone device.
 I should have not required to use a tablet of Cell phone to make it work. Amazon
 needs to work on the technology on device.
 530 Has a
 very good Bluetooth speakers sound quality is good but otherwise she's pretty
 useless when it comes to get answering questions

	reviews.title
9	was cheap, can not run chrome stuff, returned
97	Useless screen so why pay for it?
104	Amazon Fire 7 Tablet
121	Kid's Kindle
150	Have never purchased a more frustrating Device
151	not big fan
249	Good for kids
267	terrible product,bad voice quality
368	Needs to be a stand alone device
530	Good Bluetooth speaker

```
[12]: train.sentiment.value_counts()
```

```
[12]: Positive    3694
      Neutral     158
      Negative     90
      Name: sentiment, dtype: int64
```


Class Imbalance Problem In the train dataset, we have 3,749 (~95.1%) sentiments labeled as positive, and 1,58 (~4%) sentiments labeled as Neutral and 93(~2.35%) sentiments as Negative. So, it is an imbalanced classification problem.

```
[13]: pd.DataFrame(train.name.value_counts())
```

```
[13]:
```

	name	
Amazon Echo Show Alexa-enabled Bluetooth Speaker with 7" Screen		676
All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi, 16 GB - Includes Special Offers, Magenta		628
Amazon - Echo Plus w/ Built-In Hub - Silver		483
Fire Kids Edition Tablet, 7 Display, Wi-Fi, 16 GB, Blue Kid-Proof Case		446
Brand New Amazon Kindle Fire 16gb 7" Ips Display Tablet Wifi 16 Gb Blue		340
Fire Tablet, 7 Display, Wi-Fi, 16 GB - Includes Special Offers, Black		294
Amazon Tap - Alexa-Enabled Portable Bluetooth Speaker		177
Fire Kids Edition Tablet, 7 Display, Wi-Fi, 16 GB, Green Kid-Proof Case		175
Kindle E-reader - White, 6 Glare-Free Touchscreen Display, Wi-Fi - Includes Special Offers		122
Fire HD 10 Tablet, 10.1 HD Display, Wi-Fi, 16 GB - Includes Special Offers, Silver Aluminum		82
Fire Tablet with Alexa, 7" Display, 16 GB, Magenta - with Special Offers		80
Amazon Kindle E-Reader 6" Wifi (8th Generation, 2016)		76
Amazon - Kindle Voyage - 6" - 4GB - Black		65
All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi, 32 GB - Includes Special Offers, Blue		56
All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi, 32 GB - Includes Special Offers, Black		45
Fire HD 8 Tablet with Alexa, 8" HD Display, 32 GB, Tangerine - with Special Offers		43
All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi, 16 GB - Includes Special Offers, Blue		35
All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi, 32 GB - Includes Special Offers, Magenta		35
Kindle Oasis E-reader with Leather Charging Cover - Black, 6" High-Resolution Display (300 ppi), Wi-Fi - Includes Special Offers		26
Amazon 9W PowerFast Official OEM USB Charger and Power Adapter for Fire Tablets and Kindle eReaders		20

```

Amazon - Kindle Voyage - 4GB - Wi-Fi + 3G - Black
19
Kindle Oasis E-reader with Leather Charging Cover - Merlot, 6 High-Resolution
Display (300 ppi), Wi-Fi - Includes Special Offers      17
Amazon Fire TV with 4K Ultra HD and Alexa Voice Remote (Pendant Design) |
Streaming Media Player                                2

```

```

[ ]: # name = pd.DataFrame(train.name.str.split(',').tolist()).stack().unique()
# name = pd.DataFrame(name, columns=['name'])
# name

```

```

[14]: train.brand.value_counts() , test_val.brand.value_counts()

```

```

[14]: (Amazon      3942
      Name: brand, dtype: int64, Amazon      1000
      Name: brand, dtype: int64)

```

```

[15]: train.primaryCategories.value_counts()

```

```

[15]: Electronics      2562
      Electronics,Hardware      1159
      Office Supplies,Electronics      204
      Electronics,Media      17
      Name: primaryCategories, dtype: int64

```

```

[16]: test_val.primaryCategories.value_counts()

```

```

[16]: Electronics      676
      Electronics,Hardware      276
      Office Supplies,Electronics      41
      Electronics,Media      7
      Name: primaryCategories, dtype: int64

```

```

[17]: pd.DataFrame(train.categories.value_counts())

```

```

[17]: categories
      Electronics,iPad & Tablets,All Tablets,Fire Tablets,Tablets,Computers & Tablets
628
      Computers,Amazon Echo,Virtual Assistant Speakers,Audio & Video
      Components,Electronics Features,Computer Accessories,Home & Tools,See more
      Amazon Echo Show Smart Assistant - White,Smart Home Automat...      514
      Amazon Echo,Smart Home,Networking,Home & Tools,Home Improvement,Smart Home
      Automation,Voice Assistants,Amazon Home,Amazon,Smart Hub & Kits,Digital Device 3
483
      Computers,Fire Tablets,Electronics Features,Computer Accessories,Tablets,Top
      Rated,Amazon Tablets,Electronics,Kids' Tablets,iPad & Tablets,Cases &
      Bags,Electronics, Tech Toys, Movies, Music,Compute...      446

```

Computers/Tablets & Networking,Tablets & eBook Readers,Computers & Tablets,Tablets,All Tablets
340

Fire Tablets,Computers/Tablets & Networking,Tablets,All Tablets,Amazon Tablets,Frys,Computers & Tablets,Tablets & eBook Readers
294

Fire Tablets,Tablets,All Tablets,Amazon Tablets,Computers & Tablets
231

Amazon Echo,Home Theater & Audio,MP3 MP4 Player Accessories,Electronics,Portable Audio,Compact Radios Stereos,Smart Hubs & Wireless Routers,Featured Brands,Smart Home & Connected Living,Home Securi... 177

Amazon Echo,Virtual Assistant Speakers,Electronics Features,Home & Tools,Smart Home Automation,TVs Entertainment,Speakers,Smart Hub & Kits,Digital Device 3,Wireless Speakers,Smart Home,Home Improve... 162

Office,eBook Readers,Electronics Features,Walmart for Business,Tablets,Electronics,Amazon Ereaders,Office Electronics,iPad & Tablets,Kindle E-readers,All Tablets,Amazon Book Reader,Computers & Tablets
122

eBook Readers,Fire Tablets,Electronics Features,Tablets,Amazon Tablets,College Ipad & Tablets,Electronics,Electronics Deals,College Electronics,Featured Brands,All Tablets,Computers & Tablets,Back... 82

Tablets,Fire Tablets,Electronics,iPad & Tablets,Android Tablets,Computers & Tablets,All Tablets
80

Computers,Electronics Features,Tablets,Electronics,iPad & Tablets,Kindle E-readers,iPad Accessories,Used:Tablets,E-Readers,E-Readers & Accessories,Computers/Tablets & Networking,Used:Computers Acce... 76

eBook Readers,Electronics Features,Walmart for Business,Tablets,See more Amazon Kindle Voyage (Wi-Fi),Electronics,Office Electronics,iPad & Tablets,Kindle E-readers,E-Readers & Accessories,All Tabl... 65

Fire Tablets,Tablets,Computers/Tablets & Networking,Other Computers & Networking,Computers & Tablets,All Tablets
45

Tablets,Fire Tablets,Computers & Tablets,All Tablets
43

Fire Tablets,Tablets,All Tablets,Amazon Tablets
35

Tablets,Fire Tablets,Electronics,Computers,Computer Components,Hard Drives & Storage,Computers & Tablets,All Tablets
35

Kindle E-readers,Electronics Features,Computers & Tablets,E-Readers & Accessories,E-Readers,eBook Readers
26

Computers & Accessories,Tablet & E-Reader Accessories,Amazon Devices & Accessories,Electronics,Power Adapters & Cables,Computers Features,Cell Phone Accessories,Cell Phone Batteries & Power,Digital... 20

Computers & Tablets,E-Readers & Accessories,eBook Readers,Kindle E-readers

19

eBook Readers,E-Readers & Accessories,Amazon Book Reader,Computers &
Tablets,Amazon Ereaders,Kindle E-readers,E-Readers

17

Amazon SMP,TV, Video & Home Audio,Electronics,Electronics Deals,TVs
Entertainment,Digital Device 4,Tvs & Home Theater,Featured Brands,Video Devices
& TV Tuners,Consumer Electronics,TV & Video,Inter... 2

```
[ ]: # categories = pd.DataFrame(train.categories.str.split(',').tolist()).stack().  
    ↪unique()  
# categories = pd.DataFrame(categories,columns=['Categories'])  
# categories
```

```
[18]: train.dtypes
```

```
[18]: name          object  
brand           object  
categories      object  
primaryCategories object  
reviews.date    object  
reviews.text    object  
reviews.title   object  
sentiment       object  
dtype: object
```

1.0.4 Data Cleaning

```
[6]: del train['brand']  
del test_val['brand']  
del test['brand']  
  
train['reviews.date'] = train['reviews.date'].str.split('T').str[0]  
test_val['reviews.date'] = test_val['reviews.date'].str.split('T').str[0]  
test['reviews.date'] = test['reviews.date'].str.split('T').str[0]  
  
train['reviews_day'] = pd.to_datetime(train['reviews.date'], format='%Y-%m-%d').  
    ↪dt.day  
train['reviews_month'] = pd.to_datetime(train['reviews.date'],  
    ↪format='%Y-%m-%d').dt.month  
train['reviews_year'] = pd.to_datetime(train['reviews.date'],  
    ↪format='%Y-%m-%d').dt.year  
  
test_val['reviews_day'] = pd.to_datetime(test_val['reviews.date'],  
    ↪format='%Y-%m-%d').dt.day
```

```

test_val['reviews_month'] = pd.to_datetime(test_val['reviews.date'],
    ↪format='%Y-%m-%d').dt.month
test_val['reviews_year'] = pd.to_datetime(test_val['reviews.date'],
    ↪format='%Y-%m-%d').dt.year

test['reviews_day'] = pd.to_datetime(test['reviews.date'], format='%Y-%m-%d').
    ↪dt.day
test['reviews_month'] = pd.to_datetime(test['reviews.date'], format='%Y-%m-%d').
    ↪dt.month
test['reviews_year'] = pd.to_datetime(test['reviews.date'], format='%Y-%m-%d').
    ↪dt.year

del train['reviews.date']
del test['reviews.date']
del test_val['reviews.date']

train.head()

```

```

[6]:
name \
0 All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...
1 Amazon - Echo Plus w/ Built-In Hub - Silver
2 Amazon Echo Show Alexa-enabled Bluetooth Speak...
3 Fire HD 10 Tablet, 10.1 HD Display, Wi-Fi, 16 ...
4 Brand New Amazon Kindle Fire 16gb 7" Ips Displ...

categories \
0 Electronics,iPad & Tablets,All Tablets,Fire Ta...
1 Amazon Echo,Smart Home,Networking,Home & Tools...
2 Amazon Echo,Virtual Assistant Speakers,Electro...
3 eBook Readers,Fire Tablets,Electronics Feature...
4 Computers/Tablets & Networking,Tablets & eBook...

primaryCategories \
0 Electronics
1 Electronics,Hardware
2 Electronics,Hardware
3 Office Supplies,Electronics
4 Electronics

reviews.text \
0 Purchased on Black FridayPros - Great Price (e...
1 I purchased two Amazon in Echo Plus and two do...
2 Just an average Alexa option. Does show a few ...
3 very good product. Exactly what I wanted, and ...
4 This is the 3rd one I've purchased. I've bough...

reviews.title sentiment reviews_day reviews_month \

```

0	Powerful tablet	Positive	26	12
1	Amazon Echo Plus AWESOME	Positive	17	1
2	Average	Neutral	20	12
3	Greattttttt	Positive	4	8
4	Very durable!	Positive	23	1

	reviews_year
0	2016
1	2018
2	2017
3	2017
4	2017

```
[7]: name = list(set(list(train['name'])+list(test_val['name'])))
categories = list( set( list( train['categories']) +
    ↳list(test_val['categories'])))
primaryCategories = list(train['primaryCategories'].unique())

le_name = LabelEncoder()
le_cat = LabelEncoder()
le_pri = LabelEncoder()
le_name.fit(name)
le_cat.fit(categories)
le_pri.fit(primaryCategories)

train['name'] = le_name.transform(train.name)
train['categories'] = le_cat.transform(train.categories)
train['primaryCategories'] = le_pri.transform(train.primaryCategories)
test_val['name'] = le_name.transform(test_val.name)
test_val['categories'] = le_cat.transform(test_val.categories)
test_val['primaryCategories'] = le_pri.transform(test_val.primaryCategories)
test['name'] = le_name.transform(test.name)
test['categories'] = le_cat.transform(test.categories)
test['primaryCategories'] = le_pri.transform(test.primaryCategories)
```

```
[8]: train['reviews.title'].fillna(value=' ',inplace=True)
test_val['reviews.title'].fillna(value=' ',inplace=True)
test['reviews.title'].fillna(value=' ',inplace=True)
```

```
[9]: tok = WordPunctTokenizer()
ps = PorterStemmer()
wnl = WordNetLemmatizer()
negations_dic = {"isn't":"is not", "aren't":"are not", "wasn't":"was not",
    ↳"weren't":"were not",
    "haven't":"have not","hasn't":"has not","hadn't":"had",
    ↳not","won't":"will not",
```

```

        "wouldn't":"would not", "don't":"do not", "doesn't":"does_
↳not", "didn't":"did not",
        "can't":"can not", "couldn't":"could not", "shouldn't":"should_
↳not", "mightn't":"might not",
        "mustn't":"must not"}
neg_pattern = re.compile(r'\b(' + '|'.join(negations_dic.keys()) + r')\b')
def data_cleaner(text):
    text = text.replace(r"Ã", '')
    text = text.replace(r"Ã", '')
    text = text.replace(r',Ã', '\')
    text = text.lower()
    text = text.replace(r',Ã', '\')
    text = neg_pattern.sub(lambda x: negations_dic[x.group()], text)
    text = re.sub("[^a-zA-Z0-9\\"]", " ", text)
    word_tok=[x for x in tok.tokenize(text) if len(x) > 3]
#     word_stem = [ps.stem(i) for i in word_tok]
#     return (" ".join(word_stem).strip())
    word_lem = [wnl.lemmatize(i) for i in word_tok]
    return (" ".join(word_lem).strip())
for i in (train, test_val, test):
    i['reviews.text']=i['reviews.text'].apply(data_cleaner)
    i['reviews.title']=i['reviews.title'].apply(data_cleaner)

```


```
[ ]:
```

```
[58]: #test[['reviews.text', 'reviews.title']].head(10)
```

1.0.5 Visualization

```
[23]: train_len=train["reviews.text"].str.len()
test_len=test["reviews.text"].str.len()
plt.hist(train_len,bins=20,label="train reviews")
plt.hist(test_len,bins=20,label="test reviews")
plt.legend()
plt.xlim(0,2000)
plt.show()


```



```
[25]: #all_text = ' '.join([text for text in train['reviews.text']])
pos_text = ' '.join([text for text in train['reviews.
    ↳text'][train['sentiment']=='Positive']])
neg_text = ' '.join([text for text in train['reviews.
    ↳text'][train['sentiment']=='Negative']])
neu_text = ' '.join([text for text in train['reviews.
    ↳text'][train['sentiment']=='Neutral']])
```

```
[28]: wordcloud = WordCloud(width=1600, height=800, random_state=21,
    ↳max_font_size=180).generate(pos_text)
plt.figure(figsize=(12,10))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title(' POSITIVE REVIEWS')
plt.show()
```


POSITIVE REVIEWS



```
[29]: wordcloud = WordCloud(height=800, width=1600,
    ↪ random_state=21, max_font_size=180).generate(neg_text)
plt.figure(figsize=(12,10))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title(' NEGATIVE REVIEWS')
plt.show()
```

NEGATIVE REVIEWS



```
[30]: wordcloud = WordCloud(height=800, width=1600,
    ↪random_state=21,max_font_size=180).generate(neu_text)
plt.figure(figsize=(12,10))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('NEUTRAL REVIEWS')
plt.show()
```



```
[10]: le_senti = LabelEncoder()
train['sentiment'] = le_senti.fit_transform(train['sentiment'])
test_val['sentiment'] = le_senti.fit_transform(test_val['sentiment'])
```

TFIDF Vectorizer

```
[10]: tvec1 = TfidfVectorizer()
tvec2 = TfidfVectorizer()
tvec3 = TfidfVectorizer()
```

```
[11]: train1 = train.reset_index()
combi1 = train1.append(test_val, ignore_index=True, sort=False)
tvec1.fit(combi1['reviews.text'])
tvec_text1 = pd.DataFrame(tvec1.transform(train1['reviews.text']).toarray())
tvec_text2 = pd.DataFrame(tvec1.transform(test_val['reviews.text']).toarray())
tvec2.fit(combi1['reviews.title'])
tvec_title1 = pd.DataFrame(tvec2.transform(train1['reviews.title']).toarray())
tvec_title2 = pd.DataFrame(tvec2.transform(test_val['reviews.title']).toarray())
```

```

Train1 = pd.concat([train1.drop(['reviews.text', 'reviews.
↳title', 'sentiment', 'index'], axis=1), tvec_text1, tvec_title1], axis=1)
Test_Val1 = pd.concat([test_val.drop(['reviews.text', 'reviews.
↳title', 'sentiment'], axis=1), tvec_text2, tvec_title2], axis=1)
x_train1=Train1.values
y_train1=train['sentiment'].values
x_val1=Test_Val1.values
y_val1 = test_val['sentiment'].values

```

```

[12]: from nltk.tokenize import RegexpTokenizer
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction import text

punc = ['.', ',', '"', "'", '?', '!', ':', ';', '(', ')', '[', ']', '{', '}',
↳'}', '%']
stop_words = text.ENGLISH_STOP_WORDS.union(punc)

stemmer = SnowballStemmer('english')
tokenizer = RegexpTokenizer(r'[a-zA-Z\']+')

def tokenize(text):
    return [stemmer.stem(word) for word in tokenizer.tokenize(text.lower())]
tvec3 = TfidfVectorizer(stop_words = stop_words, tokenizer = tokenize,
↳max_features = 1000)
reviews=tvec3.fit_transform(combi1['reviews.text'])
words = tvec3.get_feature_names()

```

```

/opt/anaconda3/lib/python3.7/site-
packages/sklearn/feature_extraction/text.py:301: UserWarning: Your stop_words
may be inconsistent with your preprocessing. Tokenizing the stop words generated
tokens ['abov', 'afterward', 'alon', 'alreadi', 'alway', 'ani', 'anoth',
'anyon', 'anyth', 'anywher', 'becam', 'becaus', 'becom', 'befor', 'besid',
'cri', 'describ', 'dure', 'els', 'elsewher', 'empti', 'everi', 'everyon',
'everyth', 'everywher', 'fifti', 'forti', 'henc', 'hereaft', 'herebi', 'howev',
'hundr', 'inde', 'mani', 'meanwhil', 'moreov', 'nobodi', 'noon', 'noth',
'nowher', 'onc', 'onli', 'otherwis', 'ourselv', 'perhap', 'pleas', 'sever',
'sinc', 'sincer', 'sixti', 'someon', 'someth', 'sometim', 'somewher',
'themself', 'thenc', 'thereaft', 'therebi', 'therefor', 'togeth', 'twelv',
'twenti', 'veri', 'whatev', 'whenc', 'whenev', 'wherea', 'whereaft', 'wherebi',
'wherev', 'whi', 'yourself'] not in stop_words.
'stop_words.' % sorted(inconsistent))

```

1.0.6 Multinomial Naive Bayes

```
[100]: nb = MultinomialNB()
nb.fit(Train1.values,train1['sentiment'])
y_pred = nb.predict(Test_Val1.values)
y_val = test_val['sentiment']
print(confusion_matrix(y_true=y_val, y_pred=y_pred))
print(classification_report(y_true=y_val, y_pred=y_pred))
print(accuracy_score(y_val, y_pred)*100)
```

```
[[ 0  0 24]
 [ 0  0 39]
 [ 0  0 937]]
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	24
1	0.00	0.00	0.00	39
2	0.94	1.00	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

93.7

/opt/anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

```
'precision', 'predicted', average, warn_for)
```

/opt/anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

```
'precision', 'predicted', average, warn_for)
```

/opt/anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

```
'precision', 'predicted', average, warn_for)
```

Everything is classified as Positive because of Imbalance Class

1.0.7 Project Task: Week 2

1.0.8 Tackling Class Imbalance Problem:

```
[12]: train.sentiment.value_counts()
```

```
[12]: Positive    3694  
      Neutral     158  
      Negative     90  
      Name: sentiment, dtype: int64
```

```
[17]: count_2, count_1, count_0 = train.sentiment.value_counts()  
      class_2 = train[train.sentiment==2]  
      class_1 = train[train.sentiment==1]  
      class_0 = train[train.sentiment==0]
```

UnderSampling

```
[18]: class_2_under = class_2.sample(count_1)  
      train_under= pd.concat([class_2_under,class_1,class_0],axis=0)  
      print(train_under.shape)  
      print(train_under.sentiment.value_counts())
```

```
(406, 9)  
2    158  
1    158  
0     90  
Name: sentiment, dtype: int64
```

OverSampling

```
[19]: class_0_over = class_0.sample(count_2,replace=True)  
      class_1_over = class_1.sample(count_2,replace=True)  
      train_over = pd.concat([class_2,class_0_over,class_1_over],axis=0)  
      print(train_over.shape)  
      print(train_over.sentiment.value_counts())
```

```
(11082, 9)  
2    3694  
1    3694  
0    3694  
Name: sentiment, dtype: int64
```

```
[44]: lr= LogisticRegression(C=30, class_weight='balanced', solver='sag',  
                             multi_class='multinomial', n_jobs=6, random_state=40,  
                             verbose=1, max_iter=1000)
```


TFIDF Vectorizer for under-sampled data

```
[47]: train = train_under.reset_index(drop=True)
      combi = train.append(test_val , ignore_index=True)
      print(combi.shape)

      tvec1.fit(combi['reviews.text'])
      tvec_text1 = pd.DataFrame(tvec1.transform(train['reviews.text']).toarray())
      tvec_text2 = pd.DataFrame(tvec1.transform(test_val['reviews.text']).toarray())

      tvec2.fit(combi['reviews.title'])
      tvec_title1 = pd.DataFrame(tvec2.transform(train['reviews.title']).toarray())
      tvec_title2 = pd.DataFrame(tvec2.transform(test_val['reviews.title']).toarray())

      Train = pd.concat([train.drop(['reviews.text', 'reviews.
      ↳title', 'sentiment'],axis=1),tvec_text1, tvec_title1],axis=1)
      Test_Val = pd.concat([test_val.drop(['reviews.text', 'reviews.
      ↳title', 'sentiment'],axis=1),tvec_text2, tvec_title2],axis=1)
      x_train=Train.values
      y_train=train['sentiment']
      x_val=Test_Val.values
      y_val = test_val['sentiment']
```

(1406, 9)

Logistic Regresiion for under-sampled data

```
[46]: lr.fit(x_train,y_train)
      y_pred = lr.predict(x_val)
      print(confusion_matrix(y_true=y_val, y_pred=y_pred))
      print(classification_report(y_true=y_val, y_pred=y_pred))
      print('accuracy : ',accuracy_score(y_val, y_pred)*100)
```

[Parallel(n_jobs=6)]: Using backend ThreadingBackend with 6 concurrent workers.

max_iter reached after 24 seconds

```
[[ 10   6   8]
 [ 15   7  17]
 [314 195 428]]
```

	precision	recall	f1-score	support
0	0.03	0.42	0.06	24
1	0.03	0.18	0.06	39
2	0.94	0.46	0.62	937
micro avg	0.45	0.45	0.45	1000
macro avg	0.34	0.35	0.24	1000
weighted avg	0.89	0.45	0.58	1000

accuracy : 44.5

/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

"the coef_ did not converge", ConvergenceWarning)

[Parallel(n_jobs=6)]: Done 1 out of 1 | elapsed: 24.4s finished


```
[47]: lb = LabelBinarizer()
lb.fit(y_val)
y_val1 = lb.transform(y_val)
y_pred1 = lb.transform(y_pred)
print(roc_auc_score(y_val1, y_pred1, average='weighted'))
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(3):
    fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

lw=2
for i in range(3):
    plt.plot(fpr[i], tpr[i], lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of Logistic Regression of under-
↳sampled data')
plt.legend(loc="lower right")
plt.show()
```

0.5284636556242508

Receiver operating characteristic of Logistic Regression of under -sampled data



TFIDF Vectorizer for over-sampled data

```
[20]: train = train_over.reset_index(drop=True)

tvec1.fit(train['reviews.text'])
tvec_text1 = pd.DataFrame(tvec1.transform(train['reviews.text']).toarray())
tvec_text2 = pd.DataFrame(tvec1.transform(test_val['reviews.text']).toarray())

tvec2.fit(train['reviews.title'])
tvec_title1 = pd.DataFrame(tvec2.transform(train['reviews.title']).toarray())
tvec_title2 = pd.DataFrame(tvec2.transform(test_val['reviews.title']).toarray())

Train = pd.concat([train.drop(['reviews.text', 'reviews.
    ↳title', 'sentiment'],axis=1),tvec_text1, tvec_title1],axis=1)
Test_Val = pd.concat([test_val.drop(['reviews.text', 'reviews.
    ↳title', 'sentiment'],axis=1),tvec_text2, tvec_title2],axis=1)

Train.to_csv('Train.csv',encoding='utf-8')
Test_Val.to_csv('Test_Val.csv',encoding='utf-8')

x_train=Train.values
y_train=train['sentiment'].values
x_val=Test_Val.values
y_val = test_val['sentiment'].values
```


Logistic Regression for over-sampled data

```
[56]: lr.fit(x_train,y_train)
      y_pred = lr.predict(x_val)
      print(confusion_matrix(y_true=y_val, y_pred=y_pred))
      print(classification_report(y_true=y_val, y_pred=y_pred))
      print('accuracy : ',accuracy_score(y_val, y_pred)*100)
```

[Parallel(n_jobs=6)]: Using backend ThreadingBackend with 6 concurrent workers.

max_iter reached after 1000 seconds

```
[[ 13   3   8]
 [ 10  10  19]
 [214 171 552]]
```

	precision	recall	f1-score	support
0	0.05	0.54	0.10	24
1	0.05	0.26	0.09	39
2	0.95	0.59	0.73	937
micro avg	0.57	0.57	0.57	1000
macro avg	0.35	0.46	0.31	1000
weighted avg	0.90	0.57	0.69	1000

accuracy : 57.49999999999999

/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

"the coef_ did not converge", ConvergenceWarning)

[Parallel(n_jobs=6)]: Done 1 out of 1 | elapsed: 16.7min finished

Logistic Regression on over-sampled data is performing better than under-sampled data

```
[58]: lb = LabelBinarizer()
      lb.fit(y_val)
      y_val1 = lb.transform(y_val)
      y_pred1 = lb.transform(y_pred)
      print(roc_auc_score(y_val1, y_pred1, average='weighted'))
      fpr = dict()
      tpr = dict()
      roc_auc = dict()
      for i in range(3):
          fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
          roc_auc[i] = auc(fpr[i], tpr[i])

      lw=2
      for i in range(3):
          plt.plot(fpr[i], tpr[i], lw=lw,
                  label='ROC curve of class {0} (area = {1:0.2f})'
```

```


        '%.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(' Receiver operating characteristic for Logistic Regression of_
↪over-sampled data')
plt.legend(loc="lower right")
plt.show()

```

0.5804294901632032

Receiver operating characteristic for Logistic Regression of over-sampled data



Multinomial Naive Bayes

```

[109]: nb = MultinomialNB()
nb.fit(x_train,y_train)
y_pred = nb.predict(x_val)
print(confusion_matrix(y_true=y_val, y_pred=y_pred))
print(classification_report(y_true=y_val, y_pred=y_pred))
print(accuracy_score(y_val, y_pred)*100)
print(nb.score(x_train,y_train))
print(nb.score(x_val,y_val))

```

```

[[ 12   3   9]
 [  4  13  22]

```

```
[ 9  78 850]]
      precision    recall  f1-score   support

     0       0.48       0.50       0.49         24
     1       0.14       0.33       0.20         39
     2       0.96       0.91       0.94        937

 micro avg       0.88       0.88       0.88       1000
 macro avg       0.53       0.58       0.54       1000
weighted avg       0.92       0.88       0.90       1000
```

87.5

0.9589424291644107

0.875


```
[60]: lb = LabelBinarizer()
lb.fit(y_val)
y_val1 = lb.transform(y_val)
y_pred1 = lb.transform(y_pred)
print(roc_auc_score(y_val1, y_pred1, average='weighted'))
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(3):
    fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

lw=2
for i in range(3):
    plt.plot(fpr[i], tpr[i], lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of Multinomial Naive Bayes_
↪Classifier')
plt.legend(loc="lower right")
plt.show()
```

0.6979688244204161

Receiver operating characteristic of Multinomial Naive Bayes Classifier



RandomForestClassifier

```
[36]: rf= RandomForestClassifier(n_estimators=400,random_state=10).
      ↪fit(x_train,y_train)
      y_pred=rf.predict(x_val)
      print(confusion_matrix(y_true=y_val, y_pred=y_pred))
      print(classification_report(y_true=y_val, y_pred=y_pred))
      print('accuracy : ',accuracy_score(y_val, y_pred)*100)
      print(rf.score(x_train,y_train))
      print(rf.score(x_val,y_val))
```

```
[[ 6  0 18]
 [ 0  4 35]
 [ 0  0 937]]
```

	precision	recall	f1-score	support
0	1.00	0.25	0.40	24
1	1.00	0.10	0.19	39
2	0.95	1.00	0.97	937
micro avg	0.95	0.95	0.95	1000
macro avg	0.98	0.45	0.52	1000
weighted avg	0.95	0.95	0.93	1000

accuracy : 94.69999999999999


1.0
0.947

```
[41]: lb = LabelBinarizer()
lb.fit(y_val)
y_val1 = lb.transform(y_val)
y_pred1 = lb.transform(y_pred)
print(roc_auc_score(y_val1, y_pred1, average='weighted'))
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(3):
    fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

lw=2
for i in range(3):
    plt.plot(fpr[i], tpr[i], lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of Random Forest Classifier')
plt.legend(loc="lower right")
plt.show()
```

0.5793650793650793



XGBClassifier

```
[27]: xgb= XGBClassifier(n_estimators=1000,max_depth=6).fit(x_train,y_train)
y_pred=xgb.predict(x_val)
print(confusion_matrix(y_true=y_val, y_pred=y_pred))
print(classification_report(y_true=y_val, y_pred=y_pred))
print("accuracy : ",accuracy_score(y_val, y_pred)*100)
```

```
[[ 11   2  11]
 [  3  13  23]
 [  1   8 928]]
```

	precision	recall	f1-score	support
0	0.73	0.46	0.56	24
1	0.57	0.33	0.42	39
2	0.96	0.99	0.98	937
micro avg	0.95	0.95	0.95	1000
macro avg	0.75	0.59	0.65	1000
weighted avg	0.94	0.95	0.95	1000

```
95.19999999999999
```

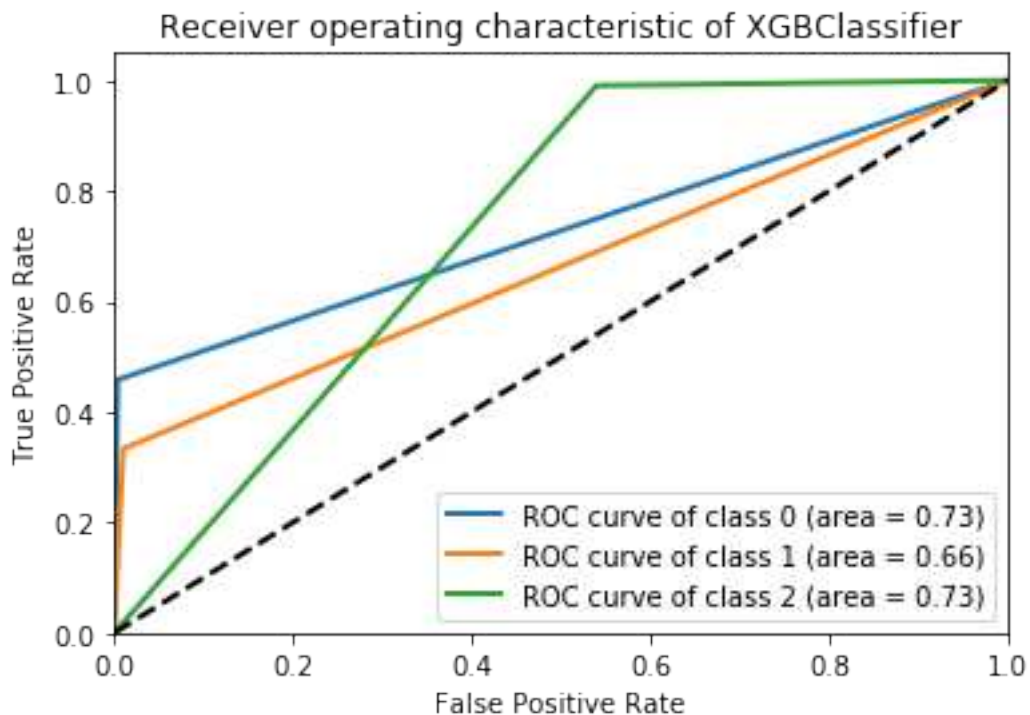
```
1.0
```

```
0.952
```

```
[40]: lb = LabelBinarizer()
lb.fit(y_val)
y_val1 = lb.transform(y_val)
y_pred1 = lb.transform(y_pred)
print(roc_auc_score(y_val1, y_pred1, average='weighted'))
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(3):
    fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

lw=2
for i in range(3):
    plt.plot(fpr[i], tpr[i], lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of XGBClassifier')
plt.legend(loc="lower right")
plt.show()
```



We can see that XGBoost is performing better in predicting all the classes.

multi-class SVM

```
[54]: svc = SVC(kernel='linear', class_weight='balanced', C=1.0, random_state=0).
      ↪ fit(x_train, y_train)
      y_pred=svc.predict(x_val)
      print(confusion_matrix(y_true=y_val, y_pred=y_pred))
      print(classification_report(y_true=y_val, y_pred=y_pred))
      print("accuracy : ",accuracy_score(y_val, y_pred)*100)
```

```
[[ 12   3   9]
 [  7  18  14]
 [ 23  82 832]]
```

	precision	recall	f1-score	support
0	0.29	0.50	0.36	24
1	0.17	0.46	0.25	39
2	0.97	0.89	0.93	937
micro avg	0.86	0.86	0.86	1000
macro avg	0.48	0.62	0.52	1000
weighted avg	0.93	0.86	0.89	1000

accuracy : 86.2

```
[55]: lb = LabelBinarizer()
      lb.fit(y_val)
      y_val1 = lb.transform(y_val)
      y_pred1 = lb.transform(y_pred)
      print(roc_auc_score(y_val1, y_pred1, average='weighted'))
      fpr = dict()
      tpr = dict()
      roc_auc = dict()
      for i in range(3):
          fpr[i], tpr[i], _ = roc_curve(y_val1[:, i], y_pred1[:, i])
          roc_auc[i] = auc(fpr[i], tpr[i])


      lw=2
      for i in range(3):
          plt.plot(fpr[i], tpr[i], lw=lw,
                  label='ROC curve of class {0} (area = {1:0.2f})'
                  ''.format(i, roc_auc[i]))

      plt.plot([0, 1], [0, 1], 'k--', lw=lw)
```



```
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of multiclass SVM Classifier')
plt.legend(loc="lower right")
plt.show()
```

0.7578666991324146



1.0.9 Project Task: Week 3

Neural Network

```
[93]: y_train2 = label_binarize(y_train1, classes=[0, 1, 2])
class_weights = class_weight.compute_class_weight('balanced',
                                                    np.unique(y_train1),
                                                    y_train1)
```

```
[87]: classifier = Sequential()
classifier.
    → add(Dense(units=100, kernel_initializer='he_uniform', activation='relu', input_dim=x_train1.
    → shape[1]))
```

```

classifier.
    →add(Dense(units=80,kernel_initializer='he_uniform',activation='relu'))
classifier.
    →add(Dense(units=80,kernel_initializer='he_uniform',activation='relu'))
classifier.add(Dense(units=3,kernel_initializer='normal',activation='softmax'))
#adam = Adam(lr=0.0001)
classifier.
    →compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
classifier.fit(x_train1,y_train2,batch_size=256,epochs=100,verbose=0)
y_pred = classifier.predict(x_val1, batch_size=256)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))

```

```

[[ 9  1 14]
 [ 0 12 27]
 [ 2  7 928]]

```

	precision	recall	f1-score	support
0	0.82	0.38	0.51	24
1	0.60	0.31	0.41	39
2	0.96	0.99	0.97	937
micro avg	0.95	0.95	0.95	1000
macro avg	0.79	0.56	0.63	1000
weighted avg	0.94	0.95	0.94	1000

```

[65]: # Using Class-Weights
classifier = Sequential()
classifier.add(Dense(units=50,activation='relu',input_dim=x_train1.shape[1]))
classifier.add(Dense(units=40,activation='relu'))
classifier.add(Dense(units=3,kernel_initializer='normal',activation='softmax'))
classifier.
    →compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
classifier.
    →fit(x_train1,y_train2,batch_size=256,epochs=100,class_weight=class_weights,verbose=0)
y_pred = classifier.predict(x_val1, batch_size=256)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))

```

```

[[ 9  2 13]
 [ 0 12 27]
 [ 2  8 927]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.82	0.38	0.51	24
1	0.55	0.31	0.39	39
2	0.96	0.99	0.97	937
micro avg	0.95	0.95	0.95	1000
macro avg	0.77	0.56	0.63	1000
weighted avg	0.94	0.95	0.94	1000

Using class-weights does not improve the performance

```
[73]: #using dropouts
classifier = Sequential()
classifier.add(Dense(units=50,activation='relu',input_dim=x_train1.shape[1]))
classifier.add(Dropout(0.2))
classifier.add(Dense(units=40,activation='relu'))
classifier.add(Dropout(0.2))
classifier.add(Dense(units=40,activation='relu'))
classifier.add(Dense(units=3,kernel_initializer='normal',activation='softmax'))
classifier.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
classifier.fit(x_train1,y_train2,batch_size=256,epochs=100,class_weight=class_weights,verbose=0)
y_pred = classifier.predict(x_val1, batch_size=256)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))
```

[[9 6 9]				
[0 15 24]				
[0 16 921]]				
	precision	recall	f1-score	support
0	1.00	0.38	0.55	24
1	0.41	0.38	0.39	39
2	0.97	0.98	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.79	0.58	0.64	1000
weighted avg	0.94	0.94	0.94	1000

Using drop out chances of predicting second class increases

```
[88]: y_train3 = label_binarize(y_train, classes=[0, 1, 2])
```

```
[90]: #for over-sampled data
classifier = Sequential()
classifier.add(Dense(units=50,activation='relu',input_dim=x_train.shape[1]))
```

```

classifier.add(Dense(units=40,activation='relu'))
classifier.add(Dense(units=150,activation='relu'))
classifier.add(Dense(units=3,kernel_initializer='normal',activation='softmax'))
classifier.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
classifier.fit(x_train,y_train3,batch_size=256,epochs=10,verbose=0)
y_pred = classifier.predict(x_val, batch_size=256)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val, y_pred_bool))
print(classification_report(y_val, y_pred_bool))

```

```

[[ 10   1  13]
 [  0  11  28]
 [  2  11 924]]

```

	precision	recall	f1-score	support
0	0.83	0.42	0.56	24
1	0.48	0.28	0.35	39
2	0.96	0.99	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.76	0.56	0.63	1000
weighted avg	0.94	0.94	0.94	1000

Using Over-sampled data for neural network does not improve the performance

1.0.10 ensemble technique using Voting Classifier: XGboost + oversampled_multinomial_NB

```

[15]: from sklearn.ensemble import VotingClassifier
model1 = MultinomialNB()
model2 = XGBClassifier(n_estimators=1000,max_depth=6)
model = VotingClassifier(estimators=[('lr', model1), ('dt', model2)],
    voting='hard')
model.fit(x_train,y_train)
y_pred = model.predict(x_val)
print(confusion_matrix(y_true=y_val, y_pred=y_pred))
print(classification_report(y_true=y_val, y_pred=y_pred))
print("accuracy : ",accuracy_score(y_val, y_pred)*100)

```

```

[[ 14   2   8]
 [  3  15  21]
 [ 14  88 835]]

```

	precision	recall	f1-score	support
0	0.45	0.58	0.51	24

1	0.14	0.38	0.21	39
2	0.97	0.89	0.93	937
micro avg	0.86	0.86	0.86	1000
macro avg	0.52	0.62	0.55	1000
weighted avg	0.92	0.86	0.89	1000

accuracy : 86.4

We can see that the above model performs almost same as oversampled multinomial model but it increases the chances of prediction of minority classes.

Sentiment Score

```
[16]: from textblob import TextBlob
def senti(x):
    return TextBlob(x).sentiment
def polarity(x):
    return TextBlob(x).polarity+1

train['senti_score'] = train['reviews.text'].apply(senti)
test_val['senti_score'] = test_val['reviews.text'].apply(senti)

train['polarity'] =train['reviews.text'].apply(polarity)
test_val['polarity'] = test_val['reviews.text'].apply(polarity)

train.senti_score.head()
```

```
[16]: 0      (0.37479166666666663, 0.6791666666666667)
1      (0.45821428571428574, 0.49821428571428567)
2      (0.69, 0.6033333333333333)
3      (0.1875, 0.4375)
4      (0.6000000000000001, 0.725)
Name: senti_score, dtype: object
```

```
[17]: Train = pd.concat([train.drop(['reviews.text', 'reviews.
→title', 'sentiment', 'senti_score'],axis=1),tvec_text1, tvec_title1],axis=1)
Test_Val = pd.concat([test_val.drop(['reviews.text', 'reviews.
→title', 'sentiment', 'senti_score'],axis=1),tvec_text2, tvec_title2],axis=1)
x_train=Train.values
y_train=train['sentiment']
x_val=Test_Val.values
y_val = test_val['sentiment']
```

```
[18]: nb = MultinomialNB()
nb.fit(x_train,y_train)
y_pred = nb.predict(x_val)
print(confusion_matrix(y_true=y_val, y_pred=y_pred))
```

```

print(classification_report(y_true=y_val, y_pred=y_pred))
print(accuracy_score(y_val, y_pred)*100)
print(nb.score(x_train,y_train))
print(nb.score(x_val,y_val))

```

```

[[ 12   4   8]
 [  3  15  21]
 [ 10  79 848]]

```

	precision	recall	f1-score	support
0	0.48	0.50	0.49	24
1	0.15	0.38	0.22	39
2	0.97	0.91	0.93	937
micro avg	0.88	0.88	0.88	1000
macro avg	0.53	0.60	0.55	1000
weighted avg	0.92	0.88	0.90	1000

```

87.5
0.9554232088070745
0.875

```

Sentiment Score does not have much affect on the performance

1.0.11 Project Task: Week 4

LSTM

```

[95]: y_train2 = label_binarize(y_train1, classes=[0, 1, 2])
epochs = 4
emb_dim = 128
batch_size = 256
model = Sequential()
model.add(Embedding(100, emb_dim, input_length=x_train1.shape[1]))
model.add(SpatialDropout1D(0.7))
model.add(LSTM(64, dropout=0.7, recurrent_dropout=0.7))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['acc'])
model.fit(x_train1, y_train2, epochs=epochs, batch_size=batch_size)
y_pred = model.predict(x_val1, batch_size=100)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))

```

```

Epoch 1/4
3942/3942 [=====] - 175s 44ms/step - loss: 0.8268 -
acc: 0.7808

```

```
Epoch 2/4
3942/3942 [=====] - 171s 43ms/step - loss: 0.3332 -
acc: 0.9371
Epoch 3/4
3942/3942 [=====] - 173s 44ms/step - loss: 0.2979 -
acc: 0.9371
Epoch 4/4
3942/3942 [=====] - 171s 43ms/step - loss: 0.2867 -
acc: 0.9371
[[ 0  0 24]
 [ 0  0 39]
 [ 0  0 937]]
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	24
1	0.00	0.00	0.00	39
2	0.94	1.00	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
[15]: #using clas_weights
y_train2 = label_binarize(y_train1, classes=[0, 1, 2])
class_weights = class_weight.compute_class_weight('balanced', np.
    ↳unique(y_train1), y_train1)
emb_dim = 128
epochs = 4
batch_size = 256
model = Sequential()
model.add(Embedding(x_train1.shape[1], emb_dim, input_length=x_train1.shape[1]))
```

```

model.add(SpatialDropout1D(0.7))
model.add(LSTM(64, dropout=0.7, recurrent_dropout=0.7))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['acc'])
model.fit(x_train1, y_train2, epochs=epochs,
    ↳batch_size=batch_size, class_weight=class_weights)
y_pred = model.predict(x_val1, batch_size=100)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))

```

WARNING:tensorflow:From /opt/anaconda3/lib/python3.7/site-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /opt/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

WARNING:tensorflow:From /opt/anaconda3/lib/python3.7/site-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/4

3942/3942 [=====] - 170s 43ms/step - loss: 0.8322 - acc: 0.8095

Epoch 2/4

3942/3942 [=====] - 165s 42ms/step - loss: 0.3274 - acc: 0.9371

Epoch 3/4

3942/3942 [=====] - 170s 43ms/step - loss: 0.3017 - acc: 0.9371

Epoch 4/4

3942/3942 [=====] - 173s 44ms/step - loss: 0.2816 - acc: 0.9371

[[0 0 24]

[0 0 39]

[0 0 937]]

precision recall f1-score support

0	0.00	0.00	0.00	24
1	0.00	0.00	0.00	39
2	0.94	1.00	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
[22]: #for over_sampled data
y_train2 = label_binarize(y_train, classes=[0, 1, 2])
emb_dim = 128
epochs = 3
batch_size = 256
model = Sequential()
model.add(Embedding(x_train.shape[1], emb_dim, input_length=x_train.shape[1]))
model.add(SpatialDropout1D(0.7))
model.add(LSTM(64, dropout=0.7, recurrent_dropout=0.7))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
↳metrics=['acc'])
model.fit(x_train, y_train2, epochs=epochs, batch_size=batch_size)
y_pred = model.predict(x_val, batch_size=100)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val, y_pred_bool))
print(classification_report(y_val, y_pred_bool))
```

Epoch 1/3

11082/11082 [=====] - 443s 40ms/step - loss: 1.1012 -
acc: 0.3352

Epoch 2/3

```

11082/11082 [=====] - 441s 40ms/step - loss: 1.1000 -
acc: 0.3302
Epoch 3/3
11082/11082 [=====] - 438s 40ms/step - loss: 1.1004 -
acc: 0.3308
[[ 0  0 24]
 [ 0  0 39]
 [ 0  0 937]]

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	24
1	0.00	0.00	0.00	39
2	0.94	1.00	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```

/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.

```

```

'precision', 'predicted', average, warn_for)

```

```

/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.

```

```

'precision', 'predicted', average, warn_for)

```

```

/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.

```

```

'precision', 'predicted', average, warn_for)

```

GRU

```

[16]: y_train2 = label_binarize(y_train1, classes=[0, 1, 2])
epochs = 3
emb_dim = 128
batch_size = 256
model = Sequential()
model.add(Embedding(x_train1.shape[1], emb_dim, input_length=x_train1.shape[1]))
#model.add(SpatialDropout1D(0.7))
model.add(GRU(64, dropout=0.3, recurrent_dropout=0.3))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
↪metrics=['acc'])

```

```

model.fit(x_train1, y_train2, epochs=epochs, batch_size=batch_size)
y_pred = model.predict(x_val1, batch_size=100)
y_pred_bool = np.argmax(y_pred, axis=1)
print(confusion_matrix(y_val1, y_pred_bool))
print(classification_report(y_val1, y_pred_bool))

```

```

Epoch 1/3
3942/3942 [=====] - 145s 37ms/step - loss: 0.7598 -
acc: 0.8595
Epoch 2/3
3942/3942 [=====] - 144s 37ms/step - loss: 0.3209 -
acc: 0.9371
Epoch 3/3
3942/3942 [=====] - 142s 36ms/step - loss: 0.2832 -
acc: 0.9371
[[ 0  0 24]
 [ 0  0 39]
 [ 0  0 937]]

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	24
1	0.00	0.00	0.00	39
2	0.94	1.00	0.97	937
micro avg	0.94	0.94	0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```

/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
'precision', 'predicted', average, warn_for)
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
'precision', 'predicted', average, warn_for)
/opt/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
'precision', 'predicted', average, warn_for)

```

[]:

We can see from above that LSTM and GPU models iare not efficient in predicting minor


classes. ANN is performing quite good in solving class imbalance problem but it cannot beat traditional ML algorithms.

Clustering of Reviews

```
[24]: print(words[250:300])
```

```
['disappoint', 'discov', 'display', 'distract', 'doe', 'doesnt', 'dollar',  
'dont', 'door', 'doorbel', 'dot', 'doubl', 'downfal', 'download', 'downsid',  
'drain', 'drawback', 'drive', 'drop', 'durabl', 'dure', 'earli', 'earlier',  
'eas', 'easi', 'easier', 'easili', 'ebook', 'echo', 'edg', 'edit', 'educ',  
'effect', 'effici', 'effort', 'electron', 'els', 'email', 'employe', 'enabl',  
'end', 'endless', 'enjoy', 'enlarg', 'entertain', 'entir', 'entri', 'environ',  
'equip', 'eread']
```

```
[33]: from sklearn.cluster import KMeans  
wcsc = []  
for i in range(1,15):  
    kmeans =   
    ↪KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0,n_jobs=-1)  
    kmeans.fit(reviews)  
    wcsc.append(kmeans.inertia_)  
plt.plot(range(1,15),wcsc)  
plt.title('The Elbow Method')  
plt.xlabel('Number of clusters')  
plt.ylabel('WCSS')  
plt.show()
```



As no proper elbow is generated, I will have to select right amount of clusters by trial and error. So, I will showcase the results of different amount of clusters to find out the right amount of clusters.

11 Clusters

```
[29]: kmeans = KMeans(n_clusters = 11, n_init = 20, n_jobs = -1)
kmeans.fit(reviews)
# We look at 6 the clusters generated by k-means.
common_words = kmeans.cluster_centers_.argsort()[:,-1:-26:-1]
for num, centroid in enumerate(common_words):
    print(str(num) + ' : ' + ', '.join(words[word] for word in centroid))
```

```
0 : veri, easi, happi, great, product, love, tablet, help, satisfi, pleas,
purchas, durabl, bought, nice, best, work, price, amazon, use, qualiti,
grandson, recommend, child, learn, enjoy
1 : echo, plus, love, alexa, amazon, great, music, sound, video, like, product,
light, devic, work, screen, famili, hous, featur, better, just, bulb, bought,
purchas, easi, thing
2 : kindl, read, love, book, great, upgrad, easi, best, light, size, like,
screen, veri, purchas, bought, better, second, model, want, batteri, origin,
replac, use, year, charg
3 : home, smart, alexa, devic, great, echo, addit, autom, control, music,
amazon, love, product, work, connect, light, purchas, video, item, easi, googl,
just, hous, abl, bulb
```

4 : gift, love, christma, bought, purchas, great, easi, wife, perfect, tablet, absolut, gave, price, product, kindl, year, kid, veri, mother, birthday, enjoy, daughter, work, good, famili

5 : great, work, product, price, easi, recommend, kid, sound, tablet, love, read, app, bought, life, friend, need, batteri, speaker, download, just, littl, book, movi, awesom, game

6 : year, love, bought, tablet, game, purchas, easi, perfect, grandson, play, great, daughter, veri, granddaught, parent, app, case, kid, warranti, christma, learn, enjoy, time, child, good

7 : like, alexa, easi, read, screen, bought, work, use, just, amazon, devic, enjoy, time, realli, music, play, book, doe, better, light, thing, need, purchas, want, product

8 : tablet, great, kid, price, app, love, amazon, need, perfect, littl, game, bought, purchas, play, like, work, child, recommend, onli, read, best, doe, want, just, time

9 : love, bought, daughter, play, game, easi, tablet, kid, alexa, grandson, christma, absolut, book, granddaught, purchas, read, great, watch, product, music, just, wife, doe, learn, screen

10 : good, tablet, price, product, veri, read, work, easi, kid, qualiti, pretti, great, sound, play, game, love, recommend, nice, size, pictur, amazon, devic, speaker, batteri, child

13 Clusters

```
[30]: kmeans = KMeans(n_clusters = 13, n_init = 20, n_jobs = -1)
kmeans.fit(reviews)
# We look at 13 the clusters generated by k-means.
common_words = kmeans.cluster_centers_.argsort()[:, -1:-26:-1]
for num, centroid in enumerate(common_words):
    print(str(num) + ' : ' + ', '.join(words[word] for word in centroid))
```

0 : alexa, music, love, home, light, smart, devic, play, question, great, turn, hous, thing, listen, speaker, control, like, amazon, just, abl, sound, news, famili, weather, kitchen

1 : game, play, love, tablet, watch, read, year, enjoy, video, book, daughter, grandson, great, bought, educ, easi, movi, learn, granddaught, download, app, realli, good, time, purchas

2 : love, bought, gift, christma, year, purchas, grandson, birthday, absolut, daughter, easi, granddaught, wife, great, tablet, parent, mother, perfect, price, gave, like, grandkid, famili, best, learn

3 : good, tablet, price, veri, product, work, qualiti, sound, easi, pretti, read, recommend, nice, great, pictur, love, devic, amazon, size, speaker, child, valu, realli, time, gift

4 : kindl, love, read, great, purchas, upgrad, better, best, model, replac, year, second, size, gift, easi, bought, veri, tablet, like, origin, screen, use, version, light, doe

5 : batteri, life, great, long, charg, easi, tablet, read, good, kindl, longer, love, light, screen, onli, veri, bought, amazon, fast, work, time, hour, better,

```

week, size
6 : like, work, easi, great, just, screen, doe, love, use, time, app, realli,
amazon, better, need, purchas, devic, bought, want, enjoy, perfect, onli, nice,
sound, size
7 : echo, plus, love, great, amazon, sound, video, music, like, alexa, home,
work, devic, product, screen, featur, famili, light, bulb, better, hous,
purchas, smart, easi, addit
8 : book, read, kindl, love, easi, great, reader, download, light, purchas,
like, want, size, perfect, just, carri, screen, need, wife, devic, game, watch,
bought, tablet, librari
9 : veri, easi, happi, love, tablet, great, purchas, bought, pleas, product,
grandson, year, help, enjoy, work, durabl, nice, satisfi, item, qualiti, price,
use, learn, friend, recommend
10 : tablet, great, price, love, app, year, need, perfect, amazon, work,
purchas, daughter, child, bought, like, littl, best, just, nice, recommend, doe,
everyth, easi, friend, time
11 : kid, great, love, tablet, easi, app, bought, good, amazon, free, price,
time, awesom, game, littl, gift, like, parent, recommend, entertain, product,
year, christma, grandson, learn
12 : great, product, work, easi, recommend, price, love, sound, best, friend,
high, gift, purchas, item, awesom, famili, qualiti, definit, veri, tablet,
devic, nice, featur, amazon, read

```

Topic Modelling

```

[13]: from sklearn.decomposition import LatentDirichletAllocation as LDA
# Helper function
def print_topics(model, count_vectorizer, n_top_words):
    words = count_vectorizer.get_feature_names()
    for topic_idx, topic in enumerate(model.components_):
        print("\nTopic #%d:" % topic_idx)
        print(" ".join([words[i]
                        for i in topic.argsort()[:n_top_words - 1:-1]]))
# Tweak the two parameters below
number_topics = 10
number_words = 10
# Create and fit the LDA model
lda = LDA(n_components=number_topics, n_jobs=-1)
lda.fit(reviews)
# Print the topics found by the LDA model
print("Topics found via LDA:")
print_topics(lda, tvec3, number_words)

```

Topics found via LDA:

Topic #0:

tablet great kindl amazon read just good app batteri book

Topic #1:

light kindl read like page screen love turn voyag button

Topic #2:

sound look great speaker easi good need love exact just

Topic #3:

parent love great control easi tablet download book purchas kid

Topic #4:

love tablet doe everyth great price awesom work bought beat

Topic #5:

recommend great good product price tablet veri easi friend high

Topic #6:

love christma gift bought kid great present tablet grandson kindl

Topic #7:

echo alexa music home love great smart light amazon devic

Topic #8:

tablet love game play year bought daughter learn granddaught easi

Topic #9:

love easi veri happi great purchas bought camera wife kindl

[]: