

Mercedes-Benz Greener Manufacturing

```
[2]: #####  
'''    Project No. 1: Mercedes-Benz Greener Manufacturing    '''  
#####  
  
# Step1: Import the required libraries  
  
# linear algebra  
import numpy as np  
# data processing, CSV file I/O (e.g. pd.read_csv)  
import pandas as pd  
# for dimensionality reduction  
from sklearn.decomposition import PCA
```

```
[3]: # Step2: Read the data from train.csv  
  
df_train = pd.read_csv('train.csv')  
# let us understand the data  
print('Size of training set: {} rows and {} columns'  
      .format(*df_train.shape))  
# print few rows and see how the data looks like  
df_train.head()
```

Size of training set: 4209 rows and 378 columns

```
[3]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 378 columns]

```
[4]: # Step3: Collect the Y values into an array

# separate the y from the data as we will use this to learn as
# the prediction output
y_train = df_train['y'].values
```

```
[5]: # Step4: Understand the data types we have

# iterate through all the columns which has X in the name of the column
cols = [c for c in df_train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))

print('Feature types:')
df_train[cols].dtypes.value_counts()
```

Number of features: 376

Feature types:

```
[5]: int64      368
      object      8
      dtype: int64
```

```
[6]: # Step5: Count the data in each of the columns

counts = [[], [], []]
for c in cols:
    typ = df_train[c].dtype
    uniq = len(np.unique(df_train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)

print('Constant features: {} Binary features: {} Categorical features: {}\n'
      .format(*[len(c) for c in counts]))
print('Constant features:', counts[0])
print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

```
[7]: # Step6: Read the test.csv data

df_test = pd.read_csv('test.csv')

# remove columns ID and Y from the data as they are not used for learning
usable_columns = list(set(df_train.columns) - set(['ID', 'y']))
y_train = df_train['y'].values
id_test = df_test['ID'].values

x_train = df_train[usable_columns]
x_test = df_test[usable_columns]
```

```
[8]: # Step7: Check for null and unique values for test and train sets

def check_missing_values(df):
    if df.isnull().any().any():
        print("There are missing values in the dataframe")
    else:
        print("There are no missing values in the dataframe")
check_missing_values(x_train)
check_missing_values(x_test)
```

There are no missing values in the dataframe

There are no missing values in the dataframe

```
[9]: # Step8: If for any column(s), the variance is equal to zero,
# then you need to remove those variable(s).
# Apply label encoder

for column in usable_columns:
    cardinality = len(np.unique(x_train[column]))
    if cardinality == 1:
        x_train.drop(column, axis=1) # Column with only one
        # value is useless so we drop it
        x_test.drop(column, axis=1)
    if cardinality > 2: # Column is categorical
        mapper = lambda x: sum([ord(digit) for digit in x])
        x_train[column] = x_train[column].apply(mapper)
        x_test[column] = x_test[column].apply(mapper)
x_train.head()
```

/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:13:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
del sys.path[0]
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:14:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
[9]:      X29  X37  X81  X229  X339  X95  X110  X138  X321  X26  ...  X241  X325  \
0      0    1    0    0    0    0    0    0    0    0  ...    0    0
1      0    1    0    1    0    0    0    0    0    0  ...    0    0
2      1    1    0    0    0    0    0    0    0    0  ...    1    0
3      1    1    0    1    0    0    0    0    0    0  ...    1    0
4      1    1    0    1    0    0    0    0    0    0  ...    0    0
```

```
      X108  X286  X263  X35  X365  X210  X297  X83
0      0    0    1    1    0    0    0    0
1      0    0    1    1    0    0    0    0
2      0    1    0    1    0    0    0    0
3      1    1    0    1    0    0    0    0
4      1    1    0    1    0    0    0    0
```

[5 rows x 376 columns]

```
[10]: # Step9: Make sure the data is now changed into numericals
```

```
print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Feature types:

```
[10]: int64    376
      dtype: int64
```

```
[11]: # Step10: Perform dimensionality reduction
      # Linear dimensionality reduction using Singular Value Decomposition of
      # the data to project it to a lower dimensional space.
n_comp = 12
pca = PCA(n_components=n_comp, random_state=420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.transform(x_test)
```

```
[12]: # Step11: Training using xgboost
```

```
import xgboost as xgb
from sklearn.metrics import r2_score
```

```

from sklearn.model_selection import train_test_split

x_train, x_valid, y_train, y_valid = train_test_split(
    pca2_results_train,
    y_train, test_size=0.2,
    random_state=4242)

d_train = xgb.DMatrix(x_train, label=y_train)
d_valid = xgb.DMatrix(x_valid, label=y_valid)
#d_test = xgb.DMatrix(x_test)
d_test = xgb.DMatrix(pca2_results_test)

params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)

watchlist = [(d_train, 'train'), (d_valid, 'valid')]

clf = xgb.train(params, d_train,
                1000, watchlist, early_stopping_rounds=50,
                feval=xgb_r2_score, maximize=True, verbose_eval=10)

```

[01:36:15] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear is now deprecated in favor of reg:squarederror.

[0] train-rmse:99.14835 valid-rmse:98.26297 train-r2:-58.35295
valid-r2:-67.63754

Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

[10] train-rmse:81.27653 valid-rmse:80.36433 train-r2:-38.88428
valid-r2:-44.91014

[20] train-rmse:66.71610 valid-rmse:65.77334 train-r2:-25.87403
valid-r2:-29.75260

[30] train-rmse:54.86915 valid-rmse:53.89103 train-r2:-17.17724
valid-r2:-19.64500

[40] train-rmse:45.24563 valid-rmse:44.22213 train-r2:-11.36018
valid-r2:-12.90149

[50] train-rmse:37.44740 valid-rmse:36.37753 train-r2:-7.46671
valid-r2:-8.40694

[60] train-rmse:31.15104 valid-rmse:30.01760 train-r2:-4.85891
valid-r2:-5.40522

[70]	train-rmse:26.08689 valid-r2:-3.40989	valid-rmse:24.90709	train-r2:-3.10881
[80]	train-rmse:22.04886 valid-r2:-2.08238	valid-rmse:20.82341	train-r2:-1.93524
[90]	train-rmse:18.85231 valid-r2:-1.20136	valid-rmse:17.59764	train-r2:-1.14586
[100]	train-rmse:16.34305 valid-r2:-0.61724	valid-rmse:15.08330	train-r2:-0.61264
[110]	train-rmse:14.40352 valid-r2:-0.22948	valid-rmse:13.15132	train-r2:-0.25259
[120]	train-rmse:12.92834 valid-r2:0.02780	valid-rmse:11.69463	train-r2:-0.00915
[130]	train-rmse:11.81810 valid-r2:0.19743	valid-rmse:10.62550	train-r2:0.15673
[140]	train-rmse:10.98451 valid-r2:0.30783	valid-rmse:9.86766	train-r2:0.27149
[150]	train-rmse:10.37958 valid-r2:0.38135	valid-rmse:9.32896	train-r2:0.34952
[160]	train-rmse:9.93296 valid-r2:0.42814	valid-rmse:8.96920	train-r2:0.40430
[170]	train-rmse:9.60073 valid-r2:0.45870	valid-rmse:8.72622	train-r2:0.44348
[180]	train-rmse:9.35625 valid-r2:0.47847	valid-rmse:8.56540	train-r2:0.47146
[190]	train-rmse:9.16724 valid-r2:0.49032	valid-rmse:8.46754	train-r2:0.49260
[200]	train-rmse:9.02452 valid-r2:0.49809	valid-rmse:8.40277	train-r2:0.50828
[210]	train-rmse:8.92179 valid-r2:0.50228	valid-rmse:8.36758	train-r2:0.51941
[220]	train-rmse:8.84667 valid-r2:0.50502	valid-rmse:8.34451	train-r2:0.52747
[230]	train-rmse:8.78950 valid-r2:0.50643	valid-rmse:8.33266	train-r2:0.53356
[240]	train-rmse:8.73954 valid-r2:0.50719	valid-rmse:8.32621	train-r2:0.53884
[250]	train-rmse:8.69151 valid-r2:0.50774	valid-rmse:8.32158	train-r2:0.54390
[260]	train-rmse:8.65760 valid-r2:0.50768	valid-rmse:8.32211	train-r2:0.54745
[270]	train-rmse:8.62205 valid-r2:0.50820	valid-rmse:8.31773	train-r2:0.55116
[280]	train-rmse:8.59481 valid-r2:0.50825	valid-rmse:8.31729	train-r2:0.55399
[290]	train-rmse:8.56563 valid-r2:0.50821	valid-rmse:8.31759	train-r2:0.55701
[300]	train-rmse:8.54286 valid-r2:0.50828	valid-rmse:8.31700	train-r2:0.55937

```

[310]   train-rmse:8.51938      valid-rmse:8.31458      train-r2:0.56178
valid-r2:0.50857
[320]   train-rmse:8.49042      valid-rmse:8.31094      train-r2:0.56476
valid-r2:0.50900
[330]   train-rmse:8.45625      valid-rmse:8.30934      train-r2:0.56825
valid-r2:0.50919
[340]   train-rmse:8.42987      valid-rmse:8.31028      train-r2:0.57094
valid-r2:0.50908
[350]   train-rmse:8.40154      valid-rmse:8.30833      train-r2:0.57382
valid-r2:0.50931
[360]   train-rmse:8.37956      valid-rmse:8.30414      train-r2:0.57605
valid-r2:0.50980
[370]   train-rmse:8.34886      valid-rmse:8.30262      train-r2:0.57915
valid-r2:0.50998
[380]   train-rmse:8.32188      valid-rmse:8.30416      train-r2:0.58187
valid-r2:0.50980
[390]   train-rmse:8.28929      valid-rmse:8.30400      train-r2:0.58514
valid-r2:0.50982
[400]   train-rmse:8.26299      valid-rmse:8.30126      train-r2:0.58776
valid-r2:0.51014
[410]   train-rmse:8.23729      valid-rmse:8.30153      train-r2:0.59032
valid-r2:0.51011
[420]   train-rmse:8.21368      valid-rmse:8.30054      train-r2:0.59267
valid-r2:0.51023
[430]   train-rmse:8.18400      valid-rmse:8.30051      train-r2:0.59561
valid-r2:0.51023
[440]   train-rmse:8.16538      valid-rmse:8.30541      train-r2:0.59745
valid-r2:0.50965
[450]   train-rmse:8.13921      valid-rmse:8.30286      train-r2:0.60002
valid-r2:0.50995
[460]   train-rmse:8.11800      valid-rmse:8.30342      train-r2:0.60211
valid-r2:0.50989
Stopping. Best iteration:
[415]   train-rmse:8.22393      valid-rmse:8.29841      train-r2:0.59165
valid-r2:0.51048

```

[14]: *# Step12: Predict your test_df values using xgboost*

```

p_test = clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index=False)

sub.head()

```

```
[14]:      ID          y
      0    1    83.118958
      1    2    97.624886
      2    3    83.798798
      3    4    77.174225
      4    5   112.502441
```

```
[15]: sub.head()
```

```
#####
'''                                '''
#####
```

```
[15]: '                                '
      End
```

```
[ ]:
```